



**A. CPU Inference Latency per Sample****B. Model Storage Footprint & Runtime RAM****Figure 10. CPU inference latency and resource footprint of FP32, INT8, and knowledge-distilled edge models.**